

# PySpark Hands-On

Mai H. Nguyen, Ph.D.

# Overview

- **Weather station measurements**
- **Data Exploration**
  - Load into Spark DataFrame
  - Describe schema
  - Show summary statistics
  - Calculate correlation between features
- **Cluster to identify different weather patterns**
  - Spark k-means
  - Parallel plots

# Get Latest from Github Repo

- **Clone/Update repo on comet**
- **If haven't cloned Summer Institute repo**
  - `git clone <URL>`
- **If already cloned Summer Institute repo**
  - `git pull <URL>`
- **<URL>**

<https://github.com/sdsc/sdsc-summer-institute-2020>

# Server Setup

- **Set up server**
  - In terminal window: `start_python_sparkr_cpu`
  - Should get something like this:  
Your notebook is here:  
`https://unkind-illicitly-mutt.comet-user-content.sdsc.edu?token=6615bbdb1a8e0fbe3ad948fb52678133`  
Submitted batch job 35032027
- **Connect to jupyter notebook**
  - In browser, paste URL of notebook from above step
- **Check queue**
  - `squeue -u $USER`

# Data Setup

- In terminal window, do the following:
- **Link to data files**  
`cd <SI2020_dir>/datasci3_scalable_machine_learning/pyspark`  
`ln -s ~/ML-data/*_weather.csv .`  
`cd ../sparkR`  
`ln -s ~/ML-data/wine*.csv .`
- **Go to spark directory**  
`cd <SI2020_dir>/datasci3_scalable_machine_learning/pyspark`

# Dataset Description

- Measurements from weather station on Mt. Woodson, San Diego
- Air temperature, humidity, wind speed, wind direction, etc.
- Three years of data: Sep. 2011 - Sep. 2014
- *minute\_weather.csv*
  - measurement every minute
- *daily\_weather.csv*
  - aggregated measurements

# Dataset Description

- Measurements from weather station on Mt. Woodson, San Diego
- Air temperature, humidity, wind speed, wind direction, etc.
- Three years of data: Sep. 2011 - Sep. 2014

- *minute\_weather.csv*

- measurement every minute


- *daily\_weather.csv*

- aggregated measurements

Clustering

Data Exploration

# Go to sparkR folder



Files

Running

Clusters

Nbextensions

Select items to perform actions on them.

☐ 0 ▾  / 1-scalableML

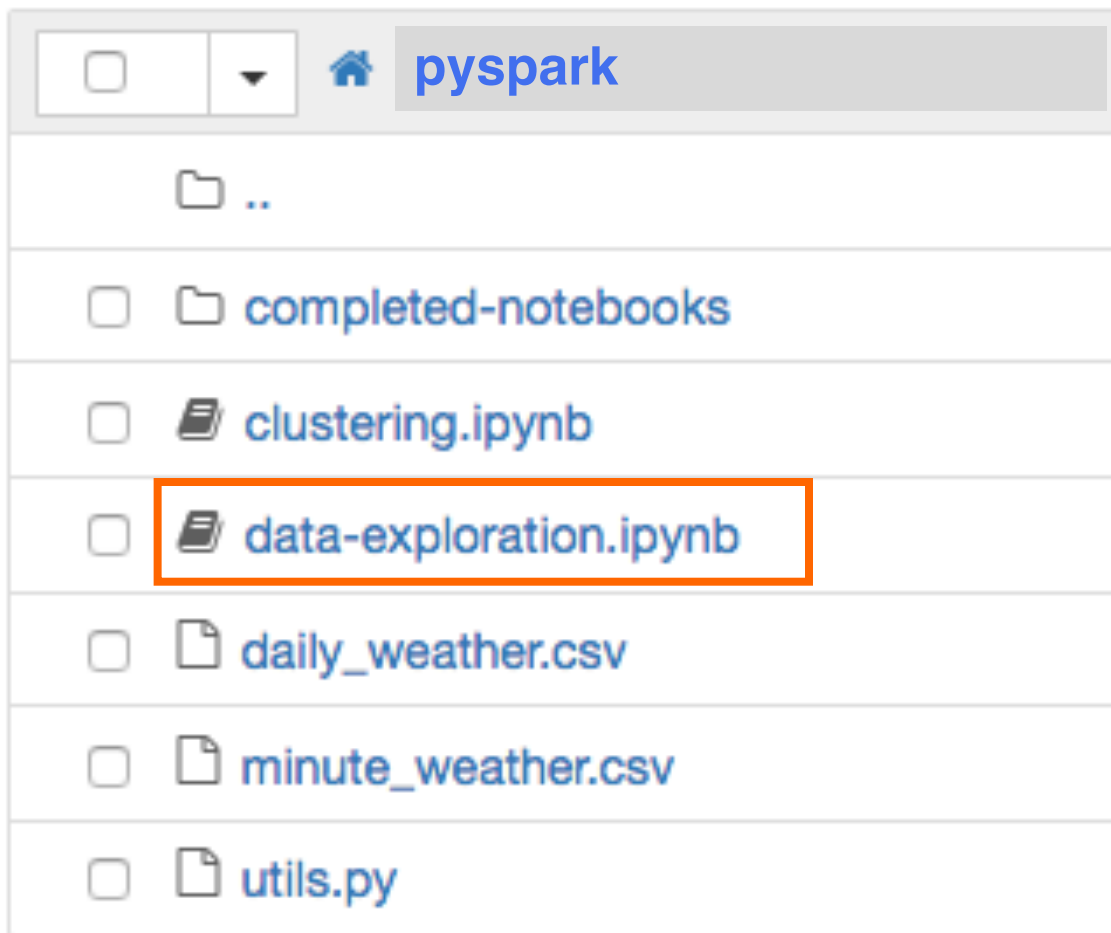
 ..

☐  pyspark

☐  sparkR



# Open Data Exploration Notebook



# Load Data into Spark DataFrame

- Start the Spark session
- Read the daily weather data into a Spark DataFrame

```
# Load data into Spark dataframe
```

```
inputfile = <<FILL-IN>>  
df = spark.read.load(inputfile, format="csv", inferSchema="true", header="true")
```

- Replace with data filename (with quotes):  
‘daily\_weather.csv’

# Examine Schema

*df.printSchema()*

```
root
|-- number: integer (nullable = true)
|-- air_pressure_9am: double (nullable = true)
|-- air_temp_9am: double (nullable = true)
|-- avg_wind_direction_9am: double (nullable = true)
|-- avg_wind_speed_9am: double (nullable = true)
|-- max_wind_direction_9am: double (nullable = true)
|-- max_wind_speed_9am: double (nullable = true)
|-- rain_accumulation_9am: double (nullable = true)
|-- rain_duration_9am: double (nullable = true)
|-- relative_humidity_9am: double (nullable = true)
|-- relative_humidity_3pm: double (nullable = true)
```

# Show Summary Statistics

*df.describe().toPandas().transpose()*

	0	1	2	3	4
summary	count	mean	stddev	min	max
number	1095	547.0	316.24357700987383	0	1094
air_pressure_9am	1092	918.8825513138094	3.184161180386833	907.9900000000024	929.3200000000012
air_temp_9am	1090	64.93300141287072	11.175514003175877	36.752000000000685	98.90599999999992
avg_wind_direction_9am	1091	142.2355107005759	69.13785928889189	15.500000000000046	343.4
avg_wind_speed_9am	1092	5.50828424225493	4.5528134655317185	0.69345139999974	23.554978199999763
max_wind_direction_9am	1092	148.95351796516923	67.23801294602953	28.89999999999991	312.19999999999993
max_wind_speed_9am	1091	7.019513529175272	5.598209170780958	1.1855782000000479	29.84077959999996
rain_accumulation_9am	1089	0.20307895225211126	1.5939521253574893	0.0	24.01999999999907
rain_duration_9am	1092	294.1080522756142	1598.0787786601481	0.0	17704.0
relative_humidity_9am	1095	34.24140205923536	25.472066802250055	6.090000000001012	92.62000000000002
relative_humidity_3pm	1095	35.34472714825898	22.524079453587273	5.3000000000006855	92.25000000000003

# Number of Rows

*df.count()*

**1095**

# First Two Rows

*df.show(2)*

```
+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+
|number|
air_pressure_9am|  air_temp_9am|avg_wind_direction_9am|avg_wind_speed_9am|max_wind
_direction_9am|max_wind_speed_9am|rain_accumulation_9am|rain_duration_9am|relative_hu
midity_9am|relative_humidity_3pm|
+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+
|  0|918.06000000000087|74.82200000000041|          271.1|
2.080354199999768|  295.39999999999986|
2.863283199999908|          0.0|          0.0|  42.42000000000046|  36.16000000000049
4|
|  1|917.3476881177097|71.40384263106537|  101.93517935618371|2.4430092157340217|
140.47154847112498|3.5333236016106238|          0.0|          0.0|  24.3286972918022
07|  19.4265967985621|
+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+
only showing top 2 rows
```

# Number and Names of Columns

*df.columns*

```
['number', 'air_pressure_9am', 'air_temp_9am',  
'avg_wind_direction_9am', 'avg_wind_speed_9am',  
'max_wind_direction_9am',  
'max_wind_speed_9am', 'rain_accumulation_9am',  
'rain_duration_9am', 'relative_humidity_9am',  
'relative_humidity_3pm']
```

*len(df.columns)*

11

# Correlation Between Air Temperature and Relative Humidity

```
df.stat.corr("air_temp_9am", "relative_humidity_9am")
```

***-0.536670...***




# Show Plots in Notebook

```
%matplotlib inline
```


# Scatter Plot of Air Temperature vs Humidity

```
df.select('air_temp_9am', 'relative_humidity_9am')  
.toPandas()  
.plot.scatter(x='air_temp_9am',  
              y='relative_humidity_9am')
```



# Histogram of Air Temperature

*df.select('air\_temp\_9am').toPandas().hist()*



# Stop Spark Session



*spark.stop()*

# Save and Exit Notebook

jupyter data-exploration Last Checkpoint: a minute ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help

New Notebook

Open...

Make a Copy...

Rename...

Save and Checkpoint

Revert to Checkpoint

Print Preview


Download as

Trusted Notebook

Close and Halt

↑ ↓ ⏮ ■ ↻ Code CellToolbar

```
lr._edgecolors == str( face ) :
```

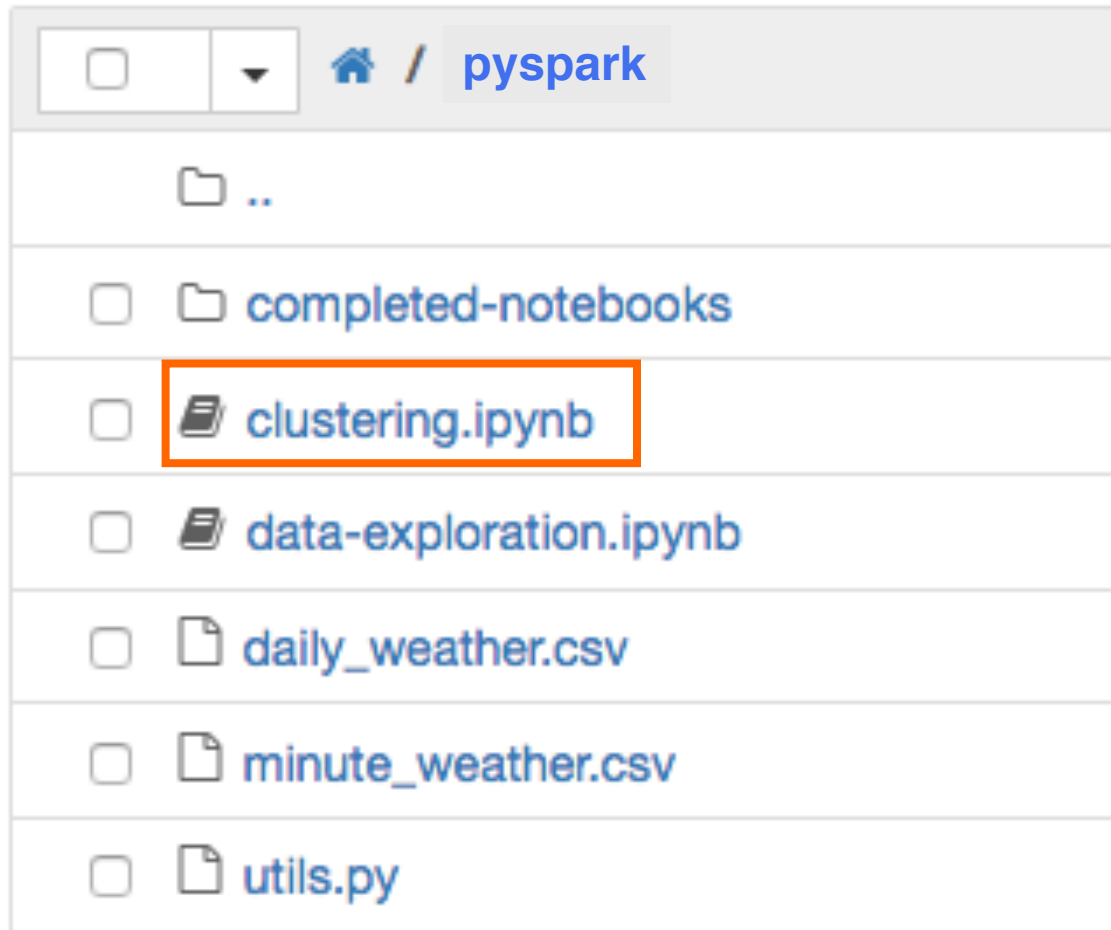


```
In [14]: # Histogram of wind speed
```

# Clustering to Identify Santa Ana Conditions

- **Strong, dry winds in Southern California**
  - wind speed  $> 30\text{mph}$
  - wind direction between 10 & 110 degrees (from east)
  - relative humidity  $< 10\%$
- **Extreme fire danger**
  - May 2014, swarm of 14 wildfires in San Diego County
  - 2008, Witch Fire,  $\sim 200,000$  acres
  - 2003, Cedar Fire,  $\sim 280,000$  acres

# Open Clustering Notebook



# Import Modules & Start Spark Session

```
# Import modules
```

```
import pyspark
from pyspark.ml.clustering import KMeans
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.feature import StandardScaler
import utils
%matplotlib inline
```

```
# Start Spark session
```

```
from pyspark.sql import SparkSession
```

```
conf = pyspark.SparkConf().setAll([('spark.master', 'local[*]'),
                                   ('spark.app.name', 'PySpark Cluster Analysis')])
spark = SparkSession.builder.config(conf=conf).getOrCreate()
```



# Load Modules & Minute Weather Data

```
# Load minute weather data
```

```
inputfile = <<FILL-IN>>
```

```
df = spark.read.csv (inputfile, inferSchema=True, header=True).cache()
```



- Replace with data filename (with quotes):
  - 'minute\_weather.csv'

# Examine Schema

*df.printSchema()*

---

```
root
|-- rowID: integer (nullable = true)
|-- hpwren_timestamp: timestamp (nullable = true)
|-- air_pressure: double (nullable = true)
|-- air_temp: double (nullable = true)
|-- avg_wind_direction: double (nullable = true)
|-- avg_wind_speed: double (nullable = true)
|-- max_wind_direction: double (nullable = true)
|-- max_wind_speed: double (nullable = true)
|-- min_wind_direction: double (nullable = true)
|-- min_wind_speed: double (nullable = true)
|-- rain_accumulation: double (nullable = true)
|-- rain_duration: double (nullable = true)
|-- relative_humidity: double (nullable = true)
```

# Count Rows and Filter Data

- **Count rows**

```
df.count()  
= 1587257
```

## **Filter data**

```
filteredDF = df.filter((df.rowID % 100) == 0)  
filteredDF.count()  
= 15873
```

# Show Summary Statistics

*filteredDF.describe().toPandas().transpose()*

	0	1	2	3	4
summary	count	mean	stddev	min	max
rowID	15873	793600.0	458228.4746717515	0	1587200
air_pressure	15873	916.8291627291587	3.0517222151797943	905.1	929.4
air_temp	15873	61.854689094688936	11.83541379082148	32.36	96.44
avg_wind_direction	15870	161.2875236294896	95.3131612965649	0.0	359.0
avg_wind_speed	15870	2.7928040327662296	2.0705061984600173	0.1	20.1
max_wind_direction	15870	162.70094517958412	92.26960112663167	0.0	359.0
max_wind_speed	15870	3.41462507876495	2.428906406812135	0.1	20.9
min_wind_direction	15870	166.64429741650915	97.82483630682509	0.0	359.0
min_wind_speed	15870	2.1522684310018896	1.7581135042599596	0.0	19.5

# Drop Samples with Null Values

```
workingDF = filteredDF.na.drop()
```

```
workingDF.count()
```

**= 15869**

# Create Feature Vector

```
featuresUsed = ['air_pressure', 'air_temp', 'avg_wind_direction', 'avg_wind_speed',  
               'max_wind_speed', 'relative_humidity']  
assembler = VectorAssembler(inputCols=featuresUsed, outputCol="features_unscaled")  
assembled = assembler.transform(workingDF)
```

# Scale Data

```
scaler = StandardScaler(inputCol="features_unscaled", outputCol="features",  
                        withStd=True, withMean=True)  
scalerModel = scaler.fit(assembled)  
scaledData = scalerModel.transform(assembled)
```

# Generate Elbow Plot

```
# Use one-third data for elbow plot
```

```
scaledData = scaledData.select("features", "rowID")
```

```
elbowset = scaledData.filter((scaledData.rowID % 3) == 0).select("features")  
elbowset.persist()  
elbowset.count()
```

5289




# Generate Clusters for Elbow Plot

```
clusters = range(5, 15)  
wsseList = utils.elbow(elbowset, clusters)
```

```
Training for cluster size 5  
.....WSSE = 9147.432309558459  
Training for cluster size 6  
.....WSSE = 8445.863736404333  
Training for cluster size 7  
.....WSSE = 8170.674504518154  
Training for cluster size 8  
.....WSSE = 7846.622384487235  
Training for cluster size 9  
.....WSSE = 7570.878097883049  
Training for cluster size 10  
.....WSSE = 7446.715915271799  
Training for cluster size 11  
.....WSSE = 7158.688288510071
```

# Show Elbow Plot

*utils.elbow\_plot(wsseList, clusters)*




# Run KMeans for $k = 9$ and Extract Cluster Centers

```
# Run KMeans for k = 9  
  
scaledDataFeat = scaledData.select("features")  
scaledDataFeat.persist()  
  
kmeans = KMeans(k=9, seed=1)  
model = kmeans.fit(scaledDataFeat)
```

```
# Extract cluster centers  
  
centers = model.clusterCenters()  
centers
```


# Cluster Capturing Humid Days

```
utils.parallel_plot(centersNamed[centersNamed['relative_humidity'] > 0.5],  
numClusters, colors=colors_used);
```




# Cluster Capturing Hot Days

```
utils.parallel_plot(centersNamed[centersNamed[air_temp'] > 0.5],  
                    numClusters, colors=colors used);
```



# Cluster Capturing Windy Days

```
utils.parallel_plot(centersNamed[centersNamed[avg_wind_speed'] > 0.5],  
numClusters, colors=colors used);
```



# Stop Spark Session

*spark.stop()*

# Clean Up

- **Exit notebook**
  - File -> Close and Halt
- **Exit Jupyter Notebook**
  - Click on 'Logout'



# References

- **Spark**
  - <https://spark.apache.org/>
- **MLlib**
  - <https://spark.apache.org/mllib/>

# Questions?

